

# **Monitoring Weak Consistency**

**Michael Emmi**  
SRI International

**Constantin Enea**  
Université Paris Diderot

# Motivation

# Concurrent Objects

## High-level abstractions

e.g. numeric & collection ADTs

## High performance

e.g. lock-free shared-memory access

## Available on modern platforms

e.g. dozens in JDK

```
/**  
 * a concurrent collection is thread-  
 * safe, but not governed by a single  
 * exclusion lock.  
 */  
package java.util.concurrent;
```

```
// we've considered these objects  
class ConcurrentHashMap { ... }  
class ConcurrentSkipListMap { ... }  
class ConcurrentSkipListSet { ... }  
class ConcurrentLinkedQueue { ... }  
class LinkedTransferQueue { ... }  
class LinkedBlockingQueue { ... }  
class ConcurrentLinkedDeque { ... }
```

```
// and there are several more
```

# Weak Consistency

## Performance optimization

synchronization vs. guarantees

## Out in the wild

e.g. collections in JDK

## Relaxed *visibility* framework

Burckhardt et al, POPL 2014

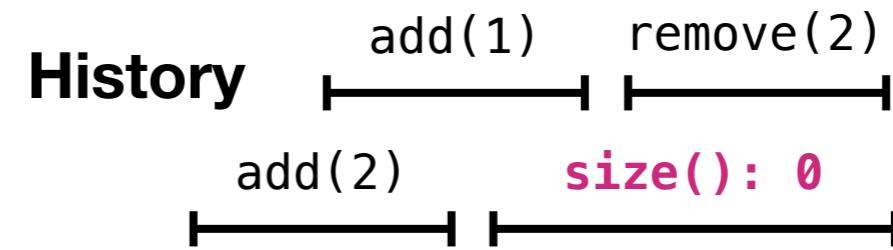
```
package java.util.concurrent;
class ConcurrentSkipListSet { ... }

/**
 * Iterators and spliterators are
 * weakly consistent...
 *
 * They are guaranteed to traverse
 * elements as they existed upon
 * construction exactly once and may
 * (but are not guaranteed to) reflect
 * any modification subsequent to
 * construction.
 */
```

# E.g. the Size Method

## History

partial order among invocations



## Linearization

total order of history

## Linearization

`add(1); add(2); remove(2); size(): 0`

## Visibility

subsequence of linearized-before  
including happens-before

## Consistency

visible sequence admitted by ADT

	visible to size	inc HB	ADT size
add(1)	✓	✓	1
add(2)	✓	✓	2
rem(2)	✓		1
			0
add(1)	✓	✓	1
add(2)	✓	✓	1
rem(2)	✓		0
			0

# **Weak Consistency Validation**

# Contributions

## Effective algorithm

consider only ***minimal*** visibilities

## Theoretical justification

minimals-only is sound

## Empirical evaluation

drastic reduction in validation time

# Algorithm

## Given history

trace of call and return actions

## Incremental search

one invocation at a time

## Enumerate linearizations

consistent with history order

## Enumerate visibilities

consistent with ADT

```
class Validator {
    h: History

    validate(lin, vis): boolean {
        // base case
        if (lin.isComplete(this.h))
            return true;

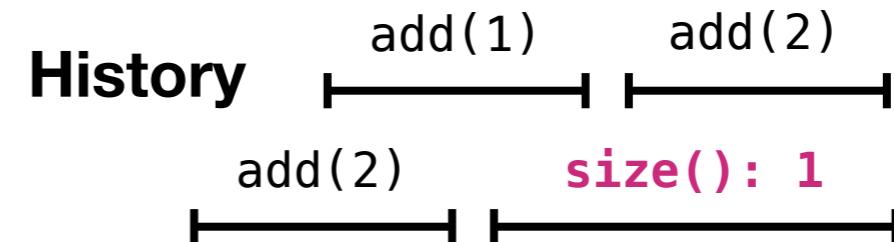
        // branch enumeration
        for (let l2 of lin.extend(this.h))
            for (let v2 of vis.extend(l2))
                if (this.validate(l2, v2))
                    return true;

        // backtracking
        return false;
    }
}

interface Linearization {
    isComplete(h: History): boolean;
    extend(h: History): this[];
}

interface Visibility {
    extend(lin: Linearization): this[];
}
```

# Minimal Visibilities



## Consistent

visible sequence admitted by ADT  
including happens-before

## Minimal

no consistent subsequence

## Theorem

sound to enumerate only minimal  
consistent visibilities

## Linearization

add(1); add(2); add(2); **size(): 1**

	visible to size		inc HB	ADT size	min
add(1)	✓	✓	✓	✓	2
add(2)		✓		✓	2
add(2)	✓		✓		2
		✓	✓	✓	1
	✓				1
		✓		✓	1
			✓		✓
				1	0

# **Experiments**

# Setup

## Java concurrent objects

reported on ConcurrentSkipListMap

## Random test generation

up to 15 invocations across 3 threads

## History recording

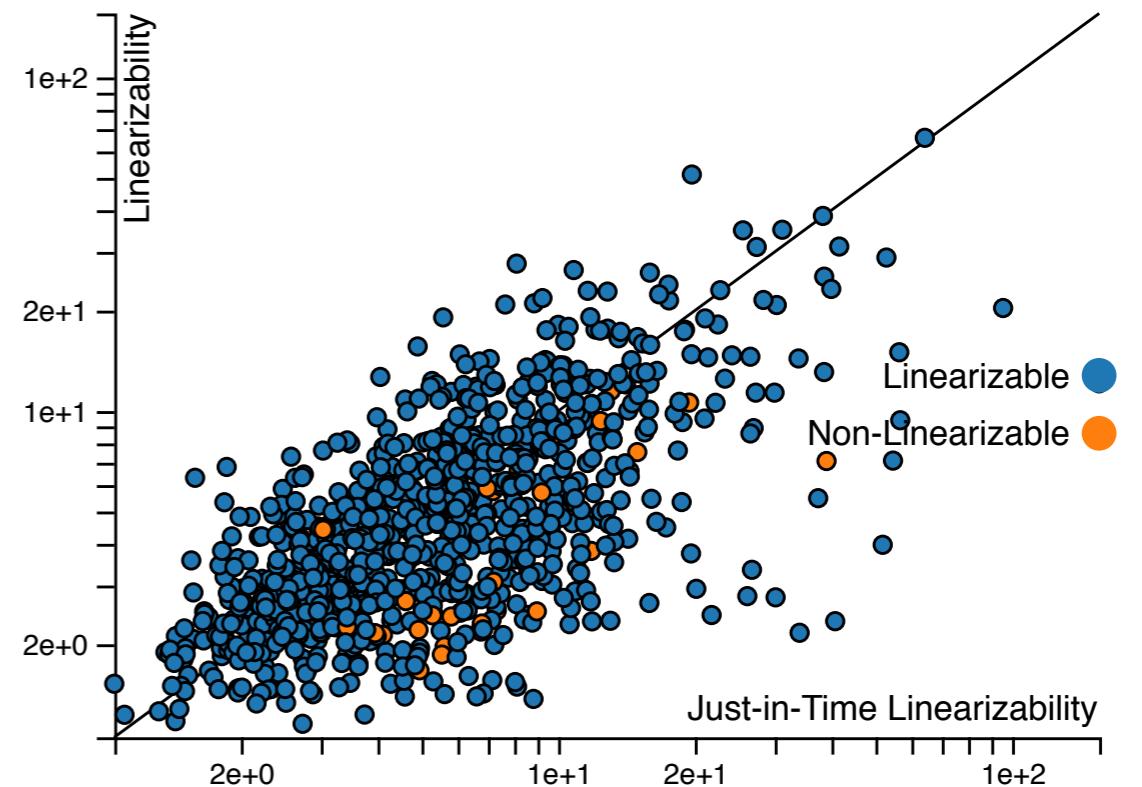
4K sampled during stress testing

## Offline, incremental validation

linearize one invocation at a time

## Basic consistency

all HB-predecessors are visible



# Baseline

**Naïve enumeration**

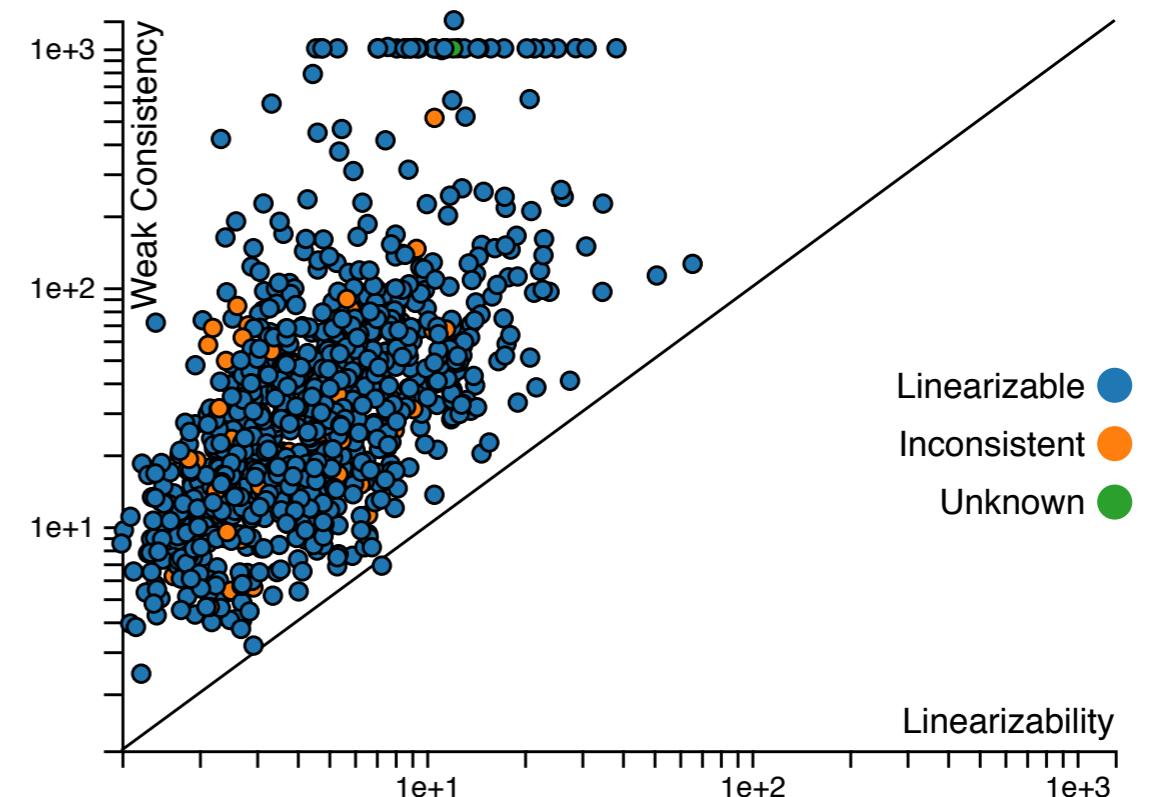
vs. linearizability

**10x overhead**

due to visibility enumeration

**Up to 3 OOM worse**

3 ms → 1000 ms (timeout)



# Optimization

**Minimals-only enumeration**

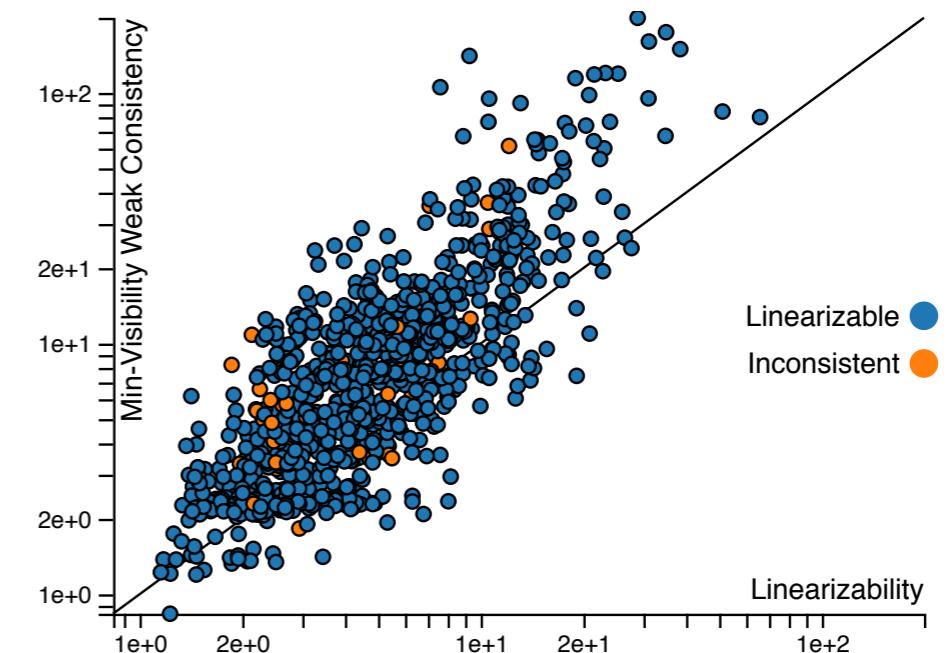
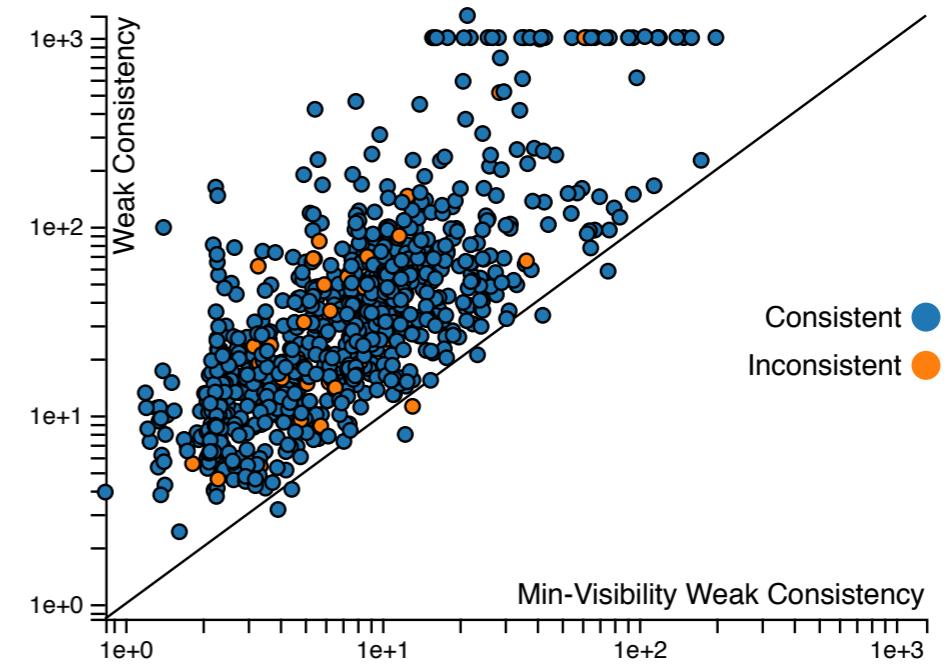
vs. naive enumeration

**Up to 2 OOM improvement**

1000 ms (timeout)  $\rightarrow$  18 ms

**2x overhead**

compared with 10x for naïve enumeration



# Related Work

## Linearization enumeration

Wing & Gong: *Testing and Verifying Concurrent Objects*. JPDC 1993

## Without specifications

Burckhardt et al: *Line-Up: a Complete and Automatic Linearizability Checker*. PLDI 2010

## With commutativity specifications

Shacham et al: *Testing Atomicity of Composed Concurrent Operations*. OOPSLA 2011

## Symbolic enumeration

Emmi, Enea, Hamza: *Monitoring Refinement via Symbolic Reasoning*. PLDI 2015

## Just-in-time enumeration

Lowe: *Testing for Linearizability*. C&C 2017